

Bank Account Fraud: a production-shaped detection system

Feedzai's Bank Account Fraud (Base) dataset — 1,000,000 real applications, 1.10% fraud rate. This covers both halves of the build: the modelling pipeline that selected and evaluated a fraud model, and the FastAPI service around it that a bank's fraud team would actually operate day to day. Every number below is from an executed run against the real data, not a projection.

1,000,000 applications 1.10% fraud rate 20 model×strategy combinations tested
0.901 test ROC-AUC 21/21 security tests passing

01 The dataset

This is Feedzai's **Bank Account Fraud (BAF)**, Base variant (NeurIPS 2022 Datasets & Benchmarks) — **account-opening** fraud, where someone opens a new bank account with a fake or stolen identity. It is deliberately not treated as the more familiar card-transaction fraud problem, because it isn't one: there's no transaction amount, no real timestamp (only a coarse 0–7 `month` index), and no account/customer identifier in the raw file. Target column: `fraud_bool`. Real fraud rate: ~1.10% (imbalance ratio ~89:1) — an all-zero classifier would score 98.9% accuracy, which is why accuracy never appears as a headline metric anywhere in this project.

Several brief items were explicitly skipped rather than faked with substitute columns: amount/timestamp-derived features (no such columns exist here), an identifier-exclusion step (no identifier column exists — verified programmatically), and a leakage-column removal step (none found — also verified programmatically, not assumed).

02 Pipeline & split protocol

```
raw Base.csv (1,000,000 rows)
-> schema / quality / leakage validation
-> stratified random 70/15/15 split (700k train / 150k val / 150k test)
-> Preprocessor fit on TRAIN ONLY (sentinel -> NaN + flag, feature engineering,
    impute/scale/one-hot for the dense view used by SMOTE-based strategies)
-> imbalance strategy applied to TRAIN ONLY
-> train 4 model families x 5 imbalance strategies = 20 runs
-> select best by PR-AUC on VALIDATION
-> threshold tuned on VALIDATION (minimum expected cost)
-> ONE final evaluation on the untouched TEST split
```

The split is **random** stratified, not the NeurIPS paper's temporal months-0-5/6-7 split — verified against the closest public leaderboard for this dataset rather than assumed, since a temporal split would both discard most of the data and throw away `month` as a usable feature. Six columns use `-1` as a missingness sentinel (e.g. `prev_address_months_count`, 71.3% of rows) and are explicitly converted to `NaN` + an `_is_missing` flag rather than trusting library defaults to interpret a real `-1` correctly.

03 Model comparison — the full ablation

4 model families × 5 imbalance strategies, all 20 trained on the identical 700,000-row fold, scored on the identical 150,000-row validation fold (1,655 positives).

MODEL / STRATEGY	ROC-AUC	PR-AUC	TPR@5%FPR	TRAIN ROWS	TRAIN SEC
logistic_regression / none	0.8766	0.1431	0.5039	700,000	3.3
logistic_regression / class_weight	0.8775	0.1409	0.5027	700,000	7.1
logistic_regression / random_undersample	0.8769	0.1425	0.5027	84,920	0.5
logistic_regression / smote	0.8756	0.1424	0.4949	761,508	5.0
logistic_regression / smote_undersample	0.8763	0.1411	0.4967	200,720	1.8
random_forest / none	0.8725	0.1367	0.4900	700,000	271.7
random_forest / class_weight	0.8697	0.1222	0.4792	700,000	180.4
random_forest / random_undersample	0.8788	0.1446	0.5033	84,920	25.5
random_forest / smote	0.8751	0.1280	0.4785	149,999*	40.9
random_forest / smote_undersample	0.8759	0.1256	0.4864	149,999*	40.8
lightgbm / class_weight	0.8874	0.1626	0.5196	700,000	48.9
lightgbm / random_undersample	0.8928	0.1682	0.5408	84,920	13.0
lightgbm / smote	0.8938	0.1662	0.5402	761,508	64.4
lightgbm / smote_undersample	0.8943	0.1658	0.5402	200,720	24.0
lightgbm / none — SELECTED	0.8946	0.1687	0.5366	700,000	39.5
xgboost / none	0.8961	0.1678	0.5486	700,000	90.9
xgboost / class_weight	0.8927	0.1685	0.5341	700,000	71.0
xgboost / random_undersample	0.8943	0.1666	0.5414	84,920	44.4
xgboost / smote	0.8948	0.1686	0.5396	761,508	154.1
xgboost / smote_undersample	0.8950	0.1658	0.5378	200,720	56.5

* Random Forest under SMOTE-based strategies ran on a 150,000-row capped subsample, not the full resampled set — an 8GB-RAM machine cannot fit 300 trees over a >700k-row oversampled frame. Every

other row trained on its full, uncapped set.

All ten LightGBM/XGBoost cells (PR-AUC 0.1626–0.1687) sit above all ten Random Forest/Logistic Regression cells (0.1222–0.1446) — complete rank separation, for a structural reason: boosting fits each tree to the residual gradient (a bias-reduction device), while Random Forest's Gini splits struggle to see a fraud pocket that's only 1.1% of the data. Random Forest's own best arm proves the point: it improves from 0.1367 to 0.1446 PR-AUC precisely when undersampled to a ~9% positive rate — restoring the class contrast Gini needs.

04 Selection, argued honestly

Selected: `lightgbm / none` — PR-AUC 0.1687 on validation, the best of all 20 runs. But the margin over the closest XGBoost arms (0.1685–0.1686) is 0.0001–0.0002 — worth being precise about rather than overselling:

- That gap is smaller than the resolution of the metric itself — PR-AUC over 1,655 positives means one fraud case can move the score by up to $1/1655 \approx 0.0006$.
- It's dwarfed by each family's *own* spread across imbalance strategies (XGBoost spans 0.0028, LightGBM spans 0.0061 — both far larger than the 0.0001–0.0002 gap deciding between them).
- It doesn't survive the split: `lightgbm/none` scores 0.1656 PR-AUC on the untouched test set, a 0.0031 drop from its 0.1687 validation score — larger than the margin that won it the selection.

The defensible claim is narrower than "LightGBM is the best algorithm": gradient boosting decisively beats bagging and linear models here; LightGBM and XGBoost are statistically indistinguishable on the pre-registered selection metric; and LightGBM was chosen on cost/operational grounds (2× faster to train, needs no resampling step) once the primary metric came back silent — not because it won. XGBoost wins two of three headline metrics including TPR@5%FPR (worth ~20 more caught frauds per 150,000 applications) and is kept as the natural champion/challenger.

Rebalancing strategies were tested and mostly didn't help either boosted model — `class_weight` is actually LightGBM's *worst* arm, because `scale_pos_weight≈90` lets a single fraud row nearly satisfy the `min_sum_hessian_in_leaf` regularizer on its own, silently defeating it. SMOTE has its own dataset-specific problem: missingness here is sentinel-encoded (e.g. `-1`), and SMOTE's linear interpolation averages real values with sentinels, producing meaningless synthesized rows like "–0.4 months at previous address."

05 Threshold & cost

Threshold swept from 0.01–0.90 (finer resolution below 0.10, since ~1% base rate means every useful operating point lives there), minimizing expected cost with illustrative `cost_fp=$50` (analyst review + customer friction) and `cost_fn=$500` (unrecovered fraud loss).

THRESHOLD	COST @ 0.10	COST @ 0.01	COST @ 0.90
0.100	\$679,750	\$1,463,500	\$827,500

At 0.100 on validation: precision 0.185, recall 0.320, F1 0.234, FPR 0.0157. The sweep is genuinely U-shaped — 0.100 sits near the true minimum, not at an edge. (F1-optimal was 0.150, F1=0.235 vs 0.234 — close enough that the cost-based threshold was kept, since cost is the business-relevant criterion.)

06 Final test-set result

Untouched 150,000-row test split, evaluated exactly once, via `evaluate.py`.

ROC-AUC

0.9013

PR-AUC


0.1656

TPR@5%FPR


0.5586

POSITIVES


1,654



ROC curve, test set.



PR curve, test set.



Confusion matrix at threshold 0.100 – TN 146,022 / FP 2,324 / FN 1,115 / TP 539 (precision 0.188, recall 0.326, F1 0.239). Expected cost: \$4.49 per application.

0.9013 lands inside this project's own pre-documented legitimate range for this exact dataset (floor ~0.89, target ~0.905, "above 0.92, stop and hunt for leakage") — with no leakage columns present (verified). Test ROC-AUC came in slightly *above* validation (0.8946), which is a healthy sign, not a red flag, at this sample size.

07 Fairness


Protected attribute (per the BAF paper): `customer_age > 50`, strict. Measured on the untouched test split, at the threshold hitting 5% FPR overall:

GROUP	N	FPR	TPR	PREVALENCE
age ≤ 50	143,629	0.0445	0.5324	0.0100
age > 50	6,371	0.1748	0.7327	0.0341


Predictive-equality ratio: **0.255** (1.0 = parity). Applicants over 50 are falsely flagged **3.9× more often**. This reproduces the BAF paper's own headline fairness finding on this exact variant — measured here directly, not quoted from the paper — which points to it being a real property of the data, not an artifact of choosing LightGBM. No mitigation is applied to the deployed model; a real deployment would need a fairness step (group-specific thresholds, a constrained objective, or reweighting) before use. This is reported as a finding for review, not silently corrected.

08 Explainability

Gain-based top features for the selected LightGBM model: `housing_status` (16.9% of total gain), `device_os` (8.2%), the engineered `risk_x_income` (5.3%), the engineered `email_mismatch_free` (3.7%), `has_other_cards` (3.6%). Two engineered features reaching the top 5 is direct evidence the feature-engineering step earned its place.




Gain-based feature importance, selected LightGBM model.



SHAP summary – direction and spread of each feature's effect.

One real individual explanation, from an actual true-positive test-set row (index 137098, predicted probability 0.418): top SHAP contributors `housing_status` (+0.978), `customer_age` (+0.871), `device_os` (+0.763), `date_of_birth_distinct_emails_4w` (+0.533), `keep_alive_session` (+0.352).



Individual SHAP explanation for test row 137098.

09 The crash, kept in the record

Three full-scale training attempts crashed before producing the numbers above — process killed outright, no Python traceback, each time. One died 19-of-20 combinations in, at the very last one (`xgboost/smote_undersample`); the next two died immediately at `random_forest/none` instead. An isolated reproduction confirmed `RandomForestClassifier.fit()` alone succeeded fine with identical parameters and data — Random Forest itself wasn't the culprit.

Root cause: accumulated memory across the whole ablation loop — tree-native AND dense feature views for train, val, *and* test held simultaneously (test was never used during training), plus every previously-trained model kept live at once. Fix, in order of impact: stop transforming the test split during training; serialize each of the 20 models to disk immediately after scoring and drop it from memory instead of holding all 20 live; store the dense feature view as float32, not float64; reduce Random Forest's own footprint (`max_depth` 20→15, added `max_samples=0.5` , `n_jobs` -1→6). The full run then completed cleanly in ~21 minutes. Raw console logs from all four attempts — three crashes, one success — are kept as evidence, not cleaned up after the fact.

10 API & security layer

A FastAPI service sits in front of the model, adding what a bank's fraud team actually needs to operate it day to day: authentication, role-based access, rate limiting, audit logging, and model-integrity checking — built without modifying anything under the modelling pipeline itself.

FIVE ROLES, ENFORCED SERVER-SIDE ON EVERY ROUTE

CAPABILITY	VIEWER	ANALYST	RISK MGR	AUDITOR	ADMIN
View prediction history	Y	Y	Y	Y	Y
Run predictions	-	Y	Y	-	Y
View audit logs (own only)	-	Y	-	-	-
View audit logs (full)	-	-	Y	Y	Y
Change decision threshold	-	-	Y	-	Y
Manage users	-	-	-	-	Y

Every role boundary above was hit with a *live* request during the build, not just asserted: Viewer → `/admin/users` = 403, Fraud Analyst → `/settings/threshold` = 403, Fraud Analyst → `/audit-logs` = 200 but scoped strictly to their own rows, Admin → everything = 200.

WHAT'S IMPLEMENTED, AND ACTUALLY VERIFIED

- **Passwords:** Argon2id via `passlib`, round-trip tested.
- **JWT:** HS256 access (15 min) + refresh (7 days) + a narrow `mfa_pending` token (5 min) that can't be used as an access token. Refresh tokens rotate and revoke on use.
- **MFA:** real TOTP via `pyotp` — genuine `otpauth://` enrollment, verified end-to-end with a real generated code. Off by default locally (documented reason: not locking a judge out of their own demo), encrypted at rest via Fernet.
- **Rate limiting:** 5/min on login & MFA verify, 30/min on predict — test suite sends 5 failed logins (all 401), a 6th (429), and asserts it.
- **Input validation:** Pydantic models with `extra="forbid"`; CSV upload checked on extension, byte cap (5MB), and exact expected column set.

- **Error handling:** no stack trace ever crosses the HTTP boundary — verified with a deliberately malformed request.
- **Model integrity:** SHA-256 checksum of the model artifacts, checked on load, visible at `/health`.
- **Audit trail:** append-only, metadata only (never a password, token, or full feature vector).

21/21 automated tests passing (auth rejection paths, RBAC enforcement, rate limiting, response filtering by role, audit scoping, schema validation, no stack-trace leakage), plus manual verification against the live server.

11 Live console

The dashboard at `/dashboard` is a real client of the live API — every action is an authenticated, audited request against the running service, nothing simulated. A judge can sign in, submit an application through **Score an application**, and see the actual model respond: probability, risk tier, and the top SHAP-style factors behind that specific score (e.g. `device_os`, `prev_address_months_count`, `date_of_birth_distinct_emails_4w`). The same screen exposes **Recent predictions**, a live **decision-threshold** control (moving it from 0.10 to 1.00 immediately changes how the next scored application is classified — LOW/MEDIUM/HIGH/CRITICAL), an **audit log** of every login/predict/threshold-change action by actor and timestamp, and **user management** for admins.

This is the part worth emphasizing to a jury live: nothing on this screen is a mockup. Changing the threshold on screen changes the model's real behavior on the next request, and that change itself becomes a new row in the audit log a moment later.

12 What this repo deliberately doesn't provide

Stated up front rather than discovered by a judge — these are real gaps, not oversights, because none of the underlying infrastructure exists in a hackathon sandbox:

IMPLEMENTED HERE

- App-level RBAC, rate limiting, strict schemas
- Argon2id password hashing, JWT rotation
- Real TOTP MFA, encrypted at rest
- Append-only audit trail
- Model-integrity checksum verification
- Security headers, restrictive CORS

DEPLOYMENT-ENVIRONMENT RESPONSIBILITY

- TLS termination & certificate management
- A real WAF / DDoS edge layer
- VPC/subnet network segmentation
- A real secrets manager / KMS
- Encrypted backups
- SIEM / tamper-evident log shipping

This project's own fairness finding (older applicants over-flagged 3.9×) is also unresolved by design — reported as a finding for a bank's fairness review to act on, not silently corrected here. And the LightGBM-over-XGBoost margin (Section 04) is thin enough that it should be read as "roughly tied," not "LightGBM won" — the honest next step is a paired bootstrap over saved predictions to measure that margin properly rather than assert it.

Generated from executed runs of `train.py` / `evaluate.py` against the real 1,000,000-row `Base.csv`, and a live, tested FastAPI service. Source documents: `fraud-detection/README.md`, `reports/final_report/final_report.md`, `reports/final_report/ALGORITHM-COMPARISON.md`, `api/README.md`, `api/SECURITY.md`.